

Parallelizing GPT-2 -“It works on my GPU”

CSCI 563

Siggy Sigler, Sam Abderholden, Dawson Matthews

Colorado School of Mines

May 4, 2026

Introduction

Serial Implementation

```
name_list: [[BOS, E, L, L, A, BOS], [name_B], [name_C]]
```

BOS = Beginning of Sequence

Serial

```
for name in name_list:
    for token in name:
        embedding(token)
        for layer in layers:
            attention(token)
            mlp(token)
            calc_loss(token)
sum_loss(tokens)
```

Serial Analysis

Big-O:

- T = sequence length, E = embedding size, L = layer, N = names
- $O(N \cdot L \cdot (C_{\text{attn}} + C_{\text{mlp}}))$
- $O(C_{\text{attn}}) = T^2 \cdot E$
- $O(C_{\text{mlp}}) = T \cdot E^2$
- Since names are mostly constant size it becomes $\approx O(N \cdot L \cdot E^2)$

Memory vs CPU Bound:

- Total Time: 6.41 s, CPU Time: 6.19
- CPU utilization $\approx \frac{6.19}{6.41} \approx 96.6\% \Rightarrow$ CPU-bound

Static Analysis:

- **99.4%** of work spent doing matrix mult, (0.4% for rmnsnorm, SM, emb)
- **71.1%** for MLP **28.3%** for ATTN
- By Amdahl's Law, with $p = 0.998$, the theoretical maximum speedup is $\frac{1}{1-p} \approx 500\times$

Serial vs. Parallel Implementation

name_list: [[BOS, E, L, L, A, BOS], [name_B], [name_C]]

BOS = Beginning of Sequence

Serial

```
for name in name_list:
    for token in name:
        embedding(token)
        for layer in layers:
            attention(token)
            mlp(token)
            calc_loss(token)
sum_loss(tokens)
```

Parallel (per batch)

```
calculate_embeddings(name_list)

for layer in layers:
    calculate_attention_scores(
        name_list)
    calculate_next_hidden_state(
        name_list)

calc_loss(name_list)
```

Parallelization Strategy

- **Batching**
 - For each sequence launch a transformer → run transformer on all sequences at once
- Store all of our information in a flattened array
 - Allows assignment of work easily by batch to threads
- Store a “workspace” on the GPU that holds weights, hidden states, etc.
- Large portions of the forward pass are parallelizable
 - **Embedding calculations**
 - **RMS normalization and linear calculations**
 - **Attention**
- **Layers must be performed sequentially**

Performance Evaluation

	Python Serial	C++ Serial	CUDA Parallel
Names: Layers: Embeddings:			
Names: Layers: Embeddings:			
Names: Layers: Embeddings:			